

Advanced Topics in Numerical Analysis: Convex and Nonsmooth Optimization Problem Set 8

Abhijit Chowdhary
ac6361@nyu.edu

New York University — April 7, 2020

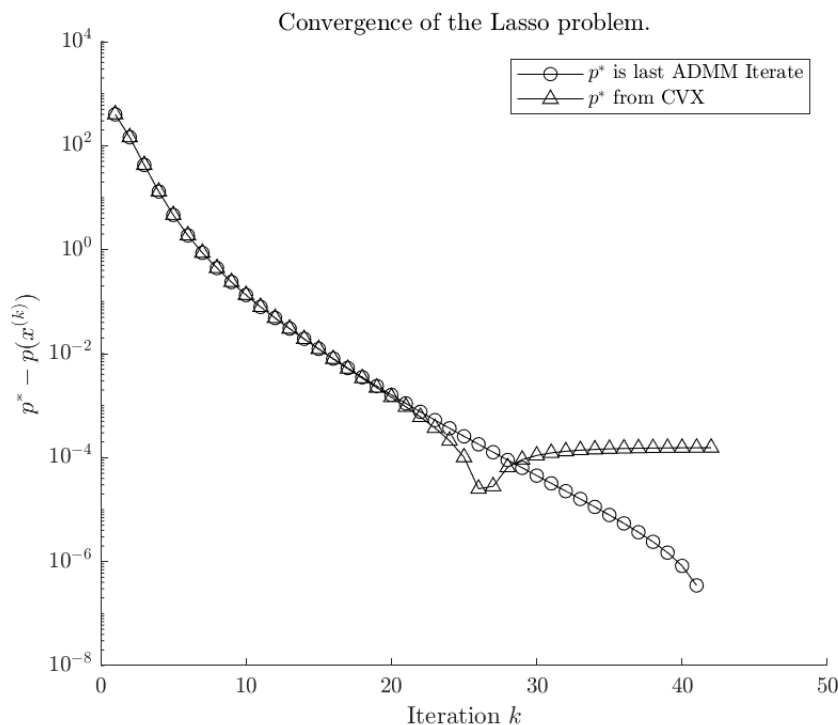


Figure 1: Log plot of the convergence of ADMM, with $\lambda = 1$ and $\rho = 1$. Note, I include two difference scenarios for p^* : (a) the last ADMM iterate and (b) the CVX solution. The dip upwards in the CVX errors is when the ADMM solution produces a smaller solution; and thus we see ADMM defeating CVX in accuracy, and by *far* in runtime (1.57s vs .06s).

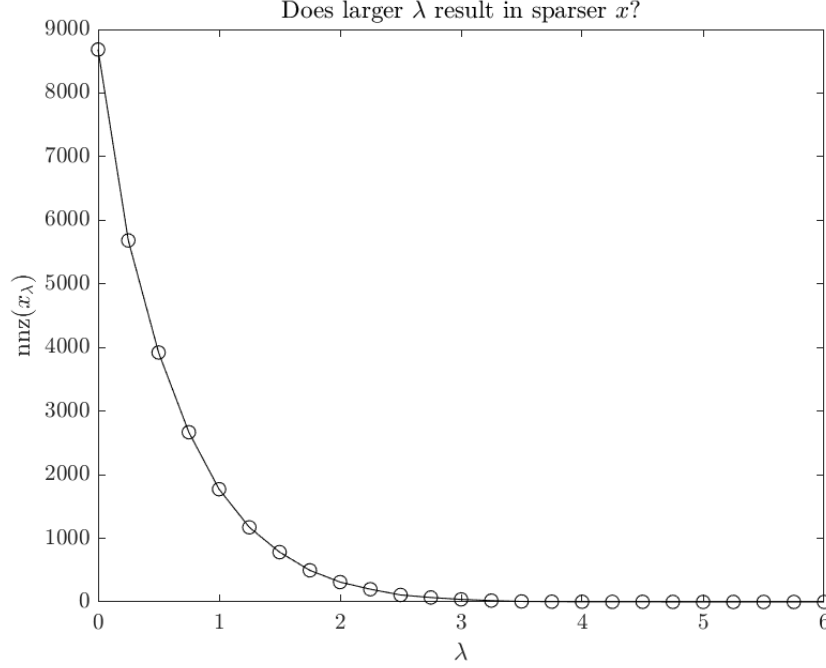


Figure 2: Indeed, as we suspect, increasing λ makes the vectors become sparser. What was surprising to me was how small λ could be before the algorithm decided that the penalty of $\|b\|_2^2/2$ wasn't difficulty to bear. Note: Initializing everything to the zero vector was important, as otherwise you have to check for values within ε of 0, rather than just running $\text{nnz}(x)$.

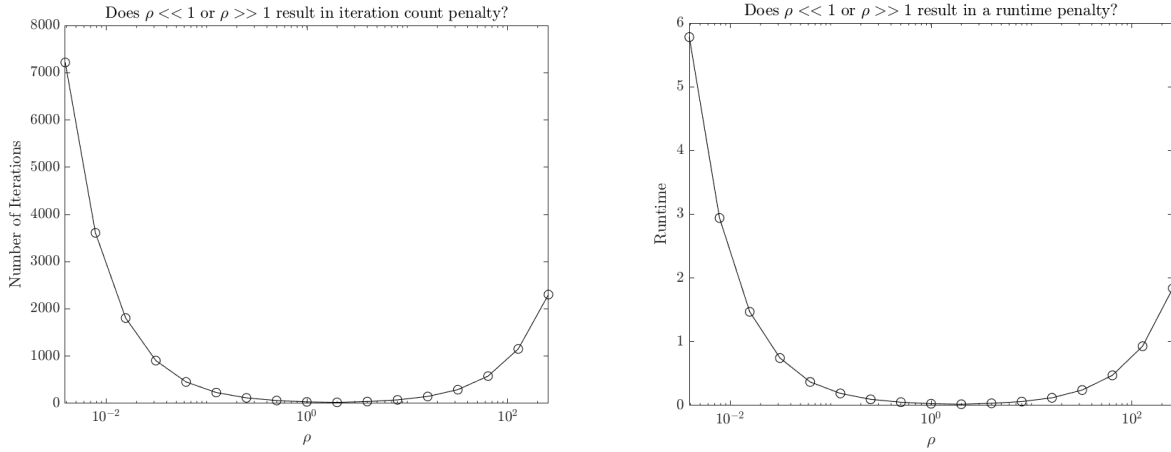


Figure 3: Interesting results here; making ρ too large or too small significantly impacts the number of iterations, and therefore the runtime of ADMM. Beautiful graphs of runtime/iteration vs ρ are produced above; they mirror each other as expected. This is likely due to the factor $\|\rho u\|$ in the computation of $\varepsilon_{\text{dual}}$.




Figure 4: If it weren't for the `nnz` function telling me that A is in fact sparse, it would be hard to tell by the graph. Interestingly, although L is sparse, it doesn't have a special pattern; just lower triangular filled in (not densely).

```

1  function [x_star, p_star, p_interm] = lasso_admm(A, b, lambda, rho, maxit, tol)
2  %% lasso_admm.m
3  %% Description:
4  % Solve the Lasso problem:
5  % Minimize (1/2) || Ax - b ||_2^2 + \lambda ||x||_1
6  % using the ADMM method.
7  %% Input:
8  % A: M x N matrix.
9  % b: N x 1 matrix.
10 % lambda: real number > 0.
11 % rho: real > 0, Augmented Lagrangian penalty parameter.
12 % Controls differentiability of augmented Lagrangian dual function.
13 % maxit: Maximum number of ADMM Iterations.
14 % tol: 1x2 matrix where:
15 % tol(1) = absolute tolerance for ADMM stopping criteria.
16 % tol(2) = relative tolerance for ADMM stopping criteria (~10^-4).
17 %% Output:
18 % x: Computed argmin to the Lasso problem.
19 % p_star: Computed minimum to the Lasso problem.
20 % p_interm: p_interm(k) computed minimum to Lasso problem at iteration k.
21
22 % Setup Lasso problem + Iteration variables
23 [M, N] = size(A);
24 S = @(a) max(a-lambda/rho, 0) - max(-a-lambda/rho, 0);
25 x = zeros(N,1); z = zeros(N,1); u = zeros(N,1); r = 0; s = 0;
26 p_interm = zeros(N,1); x_star = 0; p_star = 0;
27
28 % Precompute expensive operations
29 L = chol(A'*A + rho*speye(N), 'lower'); %try iterative solve + preconditioner?
30 ATb = A'*b;
31
32 % ADMM Iterations
33
34 for k=1:maxit
35     % ADMM Iterate + primal and dual residual computation
36     x = L' \ ( L \ ( ATb + rho*(z-u) ) );
37     z_new = S(x + u);
38     u = x + u - z_new;
39
40     r = norm(x - z_new);
41     s = norm( -rho*(z_new - z) );
42
43     z = z_new;
44

```

```

45     % Compute output values
46     p_x = (1/2)*norm(A*x-b,2)^2 + lambda*norm(x,1);
47     p_z = (1/2)*norm(A*z-b,2)^2 + lambda*norm(z,1);
48     [p_star, idx] = min([p_x p_z]); x_star = (idx == 1)*x + (idx == 2)*z;
49     p_interm(k) = p_star;
50
51     % Stopping Criteria (pg.19)
52     e_pri = sqrt(N)*tol(1) + tol(2)*max(norm(x), norm(-z)); %A, -B = I, c = 0
53     e_dual = sqrt(N)*tol(1) + tol(2)*norm(rho*u); %y = (1/rho)*y (pg.15)
54     if (r <= e_pri) && (s <= e_dual) %equation (3.12)
55         break
56     end
57 end
58 p_interm = p_interm(1:k);
59 end

```

```

1  function x = lasso_cvx(A, b, lambda)
2  %% lasso_cvx.m
3  %% Description:
4  % Solve the Lasso problem:
5  % Minimize (1/2) || Ax - b ||_2^2 + \lambda ||x||_1
6  % using CVX.
7  %% Input:
8  % A: M x N matrix.
9  % b: N x 1 matrix.
10 % lambda: real number > 0.
11 %% Output:
12 % x: Computed argmin to the Lasso problem.
13 [~,M] = size(A);
14 cvx_begin quiet
15     variable x(M)
16     minimize( 0.5*square_pos( norm(A*x-b,2) ) + lambda*norm(x,1) )
17 cvx_end
18 end

```

```

1  %% large_lasso.m
2  %% Experiments on the parameters of the ADMM method
3  clf; close all;
4  % Problem Setup
5  M = 100000; N = 10000; density = 2/M;
6  A = sprand(M,N,density); b = randn(M,1);
7
8  % Lambda experiment: Does x get sparser?
9  lambda = (0:0.25:6)';
10 nnz_lambda = zeros(size(lambda));
11 for k = 1:length(lambda)
12     l = lambda(k);
13     [x, ~, ~] = lasso_admm(A, b, l, 1, 1000, [1e-7 1e-4]);
14     fprintf('With lambda = %.2f we receive nnz(x) = %d\n', l, nnz(x));
15     nnz_lambda(k) = nnz(x);
16     %x = lasso_cvx(A, b, l);
17     %fprintf('With lambda = %.2f we receive nnz(x) = %d\n', l, N-sum(abs(x)<1e-7));
18 end %Yes, around l = 5~ we have x = 0.
19 figure();
20 plot(lambda, nnz_lambda, 'k-o');
21 title('Does larger  $\lambda$  result in sparser  $x$ ?')
22 xlabel(' $\lambda$ '); ylabel('nnz( $x_\lambda$ )');
23
24 % Rho experiment: What effect does rho have on the method?
25 lambda = 2.0;
26 rho = 2.^(-8 : 8)';
27 timings = size(length(rho));
28 for k=1:length(rho)
29     tic; [~,~,~] = lasso_admm(A,b,lambda,rho(k),10000,[1e-7 1e-4]); t = toc;
30     timings(k) = t;
31     fprintf('With rho = %.2f has runtime t=%.3f\n ', rho(k), t);
32 end %As rho << 1 or rho >> 1, the runtime suffers.
33 figure();
34 plot(rho, timings, 'k-o');
35 title('Does  $\rho << 1$  or  $\rho >> 1$  result in a runtime penalty?')
36 xlabel(' $\rho$ '); ylabel('Runtime');
37 set(gca, 'xscale', 'log');
38
39 % Rho experiment: What effect does rho have on the method?
40 num_its = size(length(rho));
41 for k=1:length(rho)
42     tic; [~,~,its] = lasso_admm(A,b,lambda,rho(k),10000,[1e-7 1e-4]); t = toc;
43     num_its(k) = length(its);
44     fprintf('With rho = %.2f numits = %d\n ', rho(k), num_its(k));

```

```

45 end %As rho << 1 or rho >> 1, the number of iterations increases suffers.
46 figure();
47 plot(rho, num_its, 'k-o');
48 title('Does  $\rho < 1$  or  $\rho > 1$  result in iteration count penalty?')
49 xlabel('$\rho$'); ylabel('Number of Iterations');
50 set(gca, 'xscale', 'log');

```

```

1  %% convergence_lasso.m
2  % Sanity checks on the convergence of ADMM
3  clf; close all;
4  % Problem Setup
5  M = 100000; N = 10000; density = 2/M;
6  A = sprand(M,N,density); b = randn(M,1);
7  lambda = 1.0;
8  obj = @(x) norm(A*x-b,2).^2/2 + lambda*norm(x,1);
9
10 x_star = lasso_cvx(A,b,lambda); p_cvxstar = obj(x_star);
11 [~,p_star,p_computed] = lasso_admm(A,b,lambda,1,100,[1e-8 1e-5]);
12
13 figure();
14 hold on;
15 plot(1:length(p_computed), p_computed-p_star, 'k-o');
16 plot(1:length(p_computed), abs(p_computed-p_cvxstar), 'k-^');
17 hold off;
18 title('Convergence of the Lasso problem. ');
19 ylabel('$p^* - p(x^{\{k\}}$'); xlabel('Iteration $k$');
20 set(gca, 'yscale', 'log');

```
